

Time Series

Ques1 – Product Demand

We would re-arrange a data little bit to get it from xls file. Attached file contains the rearranged data for both the question.



TimeSeriesAssignment.xlsx

Read the data into R

```
library(xlsx)
mydf <- read.xlsx("TimeSeriesAssignment.xlsx",header=TRUE, sheetIndex = 1)
head(mydf)
```


Visualize and understand the data

```
summary(mydf)
```

```
Min.    :15.00
1st Qu.:23.00
Median :29.50
Mean   :30.00
3rd Qu.:35.25
Max.    :53.00
```

We could see mean and median are quite close. With minimum of 15.00 and maximum of 53. Let us visualize the data

```
plot(mydf[,1], type = "l")
```



We could observe that there is a mid-year drop and then a spike at year end. Otherwise data has a general increasing pattern

Convert the data into time series object for further analysis

```
tmydf <- ts(mydf, start=c(2013, 1), end=c(2015, 12), frequency=12)
```


tmydf

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
2013	15	16	18	18	23	23	20	28	29	33	33	38
2014	23	22	28	27	31	28	22	28	32	37	34	44
2015	25	25	35	36	36	30	30	34	38	47	41	53

```
plot(tmydf)
```

```
m = lm(tmydf~time(tmydf))
```

```
abline(m)
```




We could see the trend, there is increasing mean over time

We could also see the variations across the seasons

#seasonal plots

```
seasonplot(tmydf, col = c("red", "blue", "green"), main="Seasonal Variations")
```

Seasonal Variations



This confirms the dip in July from May, and then it increases again.

On continuation from year end we could see a sharp dip in months of Jan as well which is considerably larger than mid-year dip.

a. Calculate the seasonality index using methods of averages

(Finding the index by decomposing first)

Decomposing time series means separating out the seasonality, trend and randomness from the given series

It can be done by `stl` as well as `decompose` commands

```
dmydf = decompose(tmydf, type="multiplicative")
dmydf$seasonal
plot(dmydf)
```

Decomposition of multiplicative time series



Seasonal Index

`dmydf$figure`

```
[1] 0.7849451 0.7595719 1.0046528 0.9869980 1.0411383 0.8917442 0.7662387
1.0099715 1.0739782 1.2002249 1.1266298 1.3539066
```

Let us now reconfirm this by using method of moving averages

```
mydfma= ma(tmydf,order = 12, centre = T)
mademand=tmydf/mydfma
ma_demand=t(matrix(data=mademand, nrow=12))
seasonal_demand=colMeans(ma_demand,na.rm = T)
seasonal_demand
```

```
0.7912900 0.7657117 1.0127737 0.9949761 1.0495541 0.8989524 0.7724325 1.01
81354 1.0826595 1.2099266 1.1357367 1.3648506
```

Hence we could see both are roughly same

De-seasonalize the data assuming that Y_t is product of trend and seasonality

For multiplicative time series detrended data = demand/Seasonality, therefore detrended demand can be calculated as below

```
ts_calculations <- data.frame( DetrendedData =dmydf$x/dmydf$seasonal, SeasonalIndex =
dmydf$seasonal)
ts_calculations
```

	Detrended data	Seasonality
1	19.10962	0.7849451
2	21.06450	0.7595719
3	17.91664	1.0046528
4	18.23712	0.9869980
5	22.09121	1.0411383
6	25.79215	0.8917442
7	26.10153	0.7662387
8	27.72355	1.0099715
9	27.00241	1.0739782
10	27.49485	1.2002249
11	29.29090	1.1266298
12	28.06693	1.3539066
13	29.30141	0.7849451
14	28.96368	0.7595719
15	27.87032	1.0046528
16	27.35568	0.9869980
17	29.77510	1.0411383
18	31.39914	0.8917442
19	28.71168	0.7662387
20	27.72355	1.0099715
21	29.79576	1.0739782
22	30.82756	1.2002249
23	30.17850	1.1266298
24	32.49855	1.3539066
25	31.84936	0.7849451
26	32.91328	0.7595719
27	34.83790	1.0046528
28	36.47424	0.9869980
29	34.57754	1.0411383
30	33.64194	0.8917442
31	39.15229	0.7662387
32	33.66432	1.0099715
33	35.38247	1.0739782
34	39.15933	1.2002249
35	36.39172	1.1266298
36	39.14598	1.3539066


We could also use inbuilt function `seasadj` for same
This returns seasonally adjusted data constructed by removing the seasonal component.

`seasadj(dmydf)`

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
2013	19.10962	21.06450	17.91664	18.23712	22.09121	25.79215	26.10153	27.72355
2014	27.00241	27.49485	29.30141	28.96368	27.87032	27.35568	29.77510	31.39914
2015	29.79576	30.82756	31.84936	32.91328	34.83790	36.47424	34.57754	33.64194
2016	39.15229	33.66432	35.38247	39.15933	36.39172	39.14598		

Both returns the same value

`plot(seasadj(dmydf))`



Develop the best forecasting model by comparing MAPE of MA, ES (exponential smoothing) and ARMA models. Compare the models using MAPE and Theil's coefficient.

```
#library(tseries)
```

Before we actually develop the models it is good to get some intuition about the lags.
Auto arima will anyways find this and adjust any autocorrelation

H0 (Null Hypothesis) – Series is non stationary
And Ha (Alternate Hypothesis) – Series is stationary

```
adf.test(na.omit(dmydf$random))
```

Augmented Dickey-Fuller Test

```
data: na.omit(dmydf$random)
Dickey-Fuller = -2.0431, Lag order = 2, p-value = 0.556
alternative hypothesis: stationary
```

Since p value is greater than .05 so we fail to reject the null hypothesis, indicating series to be non stationary

Let us try again but with lag of 1 to see if series becomes stationary

```
#differencing since above was not stationary
```


```
diff1 = diff(dmydf$random, differences = 1)
adf.test(na.omit(diff1))
```

Augmented Dickey-Fuller Test

```
data: na.omit(diff1)
Dickey-Fuller = -3.8435, Lag order = 2, p-value = 0.03261
alternative hypothesis: stationary
```

p value is now less than .05 so we can reject the null hypothesis and assume series is stationary.
This gives us the intuition that lag of 1 period exists in data.


```
acf(na.omit(diff1))
```



Thus now no lag is seen in residue

```
pacf(na.omit(diff1))
```


Series na.omit(diff1)



QUESTION 3: Develop the best forecasting model by comparing MAPE of MA, ES (exponential smoothing) and ARMA models. Compare the models using MAPE and Theil's coefficient.

```
library(forecast)
library(AnalyzeTS)
```

We can store actual sales data in variable, we can compare with predicted value to find accuracy of the model.

```
actual = data.frame(mydf)
```

```
# moving average model is same as built earlier
mydfma<- ma(tmydf, order=12)
plot(mydfma)
result = abs((mydfma-actual[,1])/actual[,1])*100
colMeans(result,na.rm = TRUE) # MAPE value
resid = data.frame(mydfma)
av.res(actual,resid)
```

We get MAPE value as 15.34151
And theil's coefficient as .785

We can similarly compare the two values for sma and cma

	sma	cma	mydfma	min.model
ME	2.97000	-0.10763	0.18402	cma
MAE	5.21000	4.12847	4.44444	cma
MPE	5.53498	-3.46390	-2.41424	cma
MAPE	14.71552	14.57483	15.34150	cma

MSE	47.18972	26.92693	28.81929	cma
RMSE	6.86947	5.18911	5.36836	cma
U	0.94323	0.77105	0.78479	cma

Thus out of these 3 we get best model from cma (Cumulative moving average)

Let us know see models such as arima and ets

Arima

```
auar = auto.arima(tmydf[,1])
resid = data.frame(auar$residuals)
av.res(Y=actual,E=resid)
```

	ME	MAE	MPE	MAPE	MSE	RMSE	U
auar.residuals	-0.025	1.492	-0.532	4.848	4.554	2.134	0.3367331

Exponential State Smoothing

Please refer to the following link on ets vs Holt winters

<https://robjhyndman.com/hyndsight/estimation2/>

```
s_ets = ets(tmydf)
resid = data.frame(s_ets$residuals)
av.res(Y=actual,E = resid)
```

	ME	MAE	MPE	MAPE	MSE	RMSE	U
s_ets.residuals	0.023	0.06	0.064	0.214	0.005	0.075	0.01195802

We could see that MAPE and their's coefficient turns out better for ets model among all the models

Below are two good links to know more on time series models.

<https://www.datascience.com/blog/introduction-to-forecasting-with-arima-in-r-learn-data-science-tutorials>

<https://www.r-bloggers.com/time-series-analysis-using-r-forecast-package/>

Ques2 – Television program


Before we start with solution let us again have a brief idea of data

```
summary(trp)
```

Episode		TRP	
Min.	: 1.00	Min.	: 7.230
1st Qu.:	8.25	1st Qu.:	8.633
Median	:15.50	Median	: 9.695
Mean	:15.50	Mean	: 9.512
3rd Qu.:	22.75	3rd Qu.:	10.265
Max.	:30.00	Max.	:11.990

Mean and median again are quite close and above 9.5

```
plot(trp$TRP, type = "l")
```



Develop a forecasting model using regression $Y_t = \beta_0 + \beta_1 t$, where Y_t is the TRP at time t . Is there any trend in the data? Use the regression model developed to answer?

```
tstrp = as.ts(trp)
tmod = tslm(tstrp[,2]~trend)
summary(tmod)
```

```
Call:
tslm(formula = tstrp[, 2] ~ trend)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-1.2419 -0.6874 -0.2039  0.5565  1.8002
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  8.11018    0.32775  24.745  < 2e-16 ***
trend        0.09042    0.01846   4.898 3.67e-05 ***
---

```


```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.8752 on 28 degrees of freedom
Multiple R-squared:  0.4614, Adjusted R-squared:  0.4422
F-statistic: 23.99 on 1 and 28 DF, p-value: 3.669e-05
```

Equation could be written as $Y_t = 8.11 + .0904t$

Is there an auto-correlation in the data? Conduct an appropriate hypothesis test to justify your answer.

Lag in residual could be seen from acf plot
`acf(tmod$residuals)`



Since vertical lines does not cross the confidence interval line (blue lines), hence we can conclude that no lag is there.

Let us also test this with statistical methods

Durbin-Watson test

H_0 = autocorrelation of the disturbances is 0

H_a = Autocorrelation is not 0

```
dwtest(tmod, alt="two.sided")
Durbin-watson test
```

data: tmod

DW = 1.5778, p-value = 0.1653

alternative hypothesis: true autocorrelation is not 0

Since p value is > 0.05 , thus we fail to reject the null hypothesis, which means we can reject the hypothesis that autocorrelation is 0

The television channel would like to replace the program with a new program, the average TRP of new program will be 8 points. Based on the model developed, comment whether they should replace the program with a new program.

We could see that our regression equation is $\underline{Y_t = 8.11 + .0904t}$

Also, there is an increasing trend overall for episode.

Considering these two points it would not be a good idea to replace it with series having TRP of 8
Calculate the probability that the TRP for episode 31 will be more than 10.

Prediction for next 5 episode shows the following data:

Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
31	10.91315	9.686491	12.13981	8.998754	12.82754
32	11.00357	9.769275	12.23786	9.077257	12.92988
33	11.09399	9.851632	12.33634	9.155095	13.03288
34	11.18440	9.933572	12.43524	9.232282	13.13653
35	11.27482	10.015103	12.53454	9.308830	13.24081

Mean value for TRP = 9.51

Sd can be calculated with help of confidence interval

We know that for any values in 95% confidence interval lies in 1.96 times standard deviation

Thus $12.82754 = 10.91315 + 1.96 * sd = .9767$

Now assuming normal distribution we can find the probability as

`pnorm(10,10.91315,.9767, lower.tail = FALSE) = 82.5089`

~83% of chance that probability is greater than 10